

F&B Process Anomaly Detection

Code:

```
import streamlit as st

import pandas as pd

from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier

import matplotlib.pyplot as plt

import seaborn as sns

# --- Page Configuration ---

st.set_page_config(
    page_title="Wine Quality Predictor",
    page_icon="🍷",
    layout="wide"
)

# --- Model Training ---

@st.cache_data
def train_model(df):
    """Loads data and trains the Random Forest model."""

    df['quality_category'] = df['quality'].apply(lambda x: 1 if x >= 7 else 0)

    X = df.drop(['quality', 'quality_category'], axis=1)

    y = df['quality_category']

    X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25, random_state=42, stratify=y)

    model = RandomForestClassifier(n_estimators=200, random_state=42)
```

```

    model.fit(X_train, y_train)

    return model, X

# --- Load Data and Train Model ---

try:

    # --- THIS IS THE FINAL FIX ---

    # Based on your test, we are now using the correct comma separator.

    df = pd.read_csv('winequality-red.csv', sep=',')

    model, X = train_model(df)

except FileNotFoundError:

    st.error("The 'winequality-red.csv' file was not found. Please make sure it's in the same folder as this script.")

    st.stop()

except Exception as e:

    st.error(f"An error occurred during data loading or model training: {e}")

    st.stop()

# --- Dashboard Layout ---

st.title("🍷 F&B Process Anomaly Detection")

st.header("Wine Quality Prediction Dashboard")

# --- Sidebar for User Input ---

st.sidebar.header("Input Wine Properties")

st.sidebar.markdown("Use the sliders to input the chemical properties of a wine batch.")

user_inputs = {}

for feature in X.columns:

    min_val = float(X[feature].min())

    max_val = float(X[feature].max())

```

```

mean_val = float(X[feature].mean())

user_inputs[feature] = st.sidebar.slider(

    label=f"{feature}",

    min_value=min_val,

    max_value=max_val,

    value=mean_val,

    step=0.01

)

# --- Main Panel ---

if st.sidebar.button("Predict Quality"):

    user_df = pd.DataFrame([user_inputs])

    prediction = model.predict(user_df)[0]

    prediction_proba = model.predict_proba(user_df)[0]


    st.subheader("Prediction Result")

    col1, col2 = st.columns(2)


    with col1:

        if prediction == 0: # Bad quality

            confidence = prediction_proba[0]

            st.error(f"🚨 QUALITY ALERT: Predicted Bad Quality")

            st.metric(label="Confidence", value=f"{confidence:.2%}")

            st.warning("This batch shows signs of a process anomaly.")

        else: # Good quality

            confidence = prediction_proba[1]

            st.success(f"Predicted Good Quality")

            st.metric(label="Confidence", value=f"{confidence:.2%}")

            st.info("This batch meets the quality standards.")

```

with col2:

```
st.subheader("Model Insights")
```

```
st.markdown("This chart shows which properties are most important for predicting quality.")
```

```
feature_importances = pd.Series(model.feature_importances_, index=X.columns)
```

```
feature_importances = feature_importances.sort_values(ascending=False)
```

```
fig, ax = plt.subplots()
```

```
sns.barplot(x=feature_importances, y=feature_importances.index, ax=ax)
```

```
ax.set_title('Feature Importances')
```

```
ax.set_xlabel('Importance')
```


```
st.pyplot(fig)
```


--- Instructions Section ---


```
st.markdown("<hr>", unsafe_allow_html=True)
```


```
with st.expander("  How to Use This Dashboard"):
```

```
    st.markdown("""
```

```
    -  **Adjust Sliders:** Use the sliders on the left to input your wine's properties.
```

```
    -  **Predict:** Click the 'Predict Quality' button to see the result.
```

```
    -  **View Result:** The dashboard will show the predicted quality (Good or Bad) and the model's confidence.
```

```
    -  **Get Insights:** The 'Feature Importances' chart reveals which factors most influence the prediction.
```

```
    """)
```

Dataset link: [Red wine dataset](#)

Why I chose Random Forest for this project?

1. Handles Non-Linearity Well

- The relationship between chemical properties (like acidity, alcohol, sulphates) and wine quality is not purely linear. Random Forest, being an ensemble of decision trees, can capture complex, non-linear relationships better than simple linear models.

2. Robustness Against Overfitting

- Unlike a single decision tree, Random Forest builds multiple trees and averages their results. This reduces the risk of overfitting, which is important since the dataset is not very large.

3. High Accuracy and Reliability

- Random Forest is known for strong predictive performance on classification tasks. In benchmark tests on the **Wine Quality dataset**, Random Forest often achieves higher accuracy compared to Logistic Regression or a single Decision Tree.

4. Works Well with Imbalanced Data

- The dataset has more "average" wines and fewer "good" ones. Random Forest can handle imbalance better through class weighting and bootstrapping.

5. Feature Importance Analysis

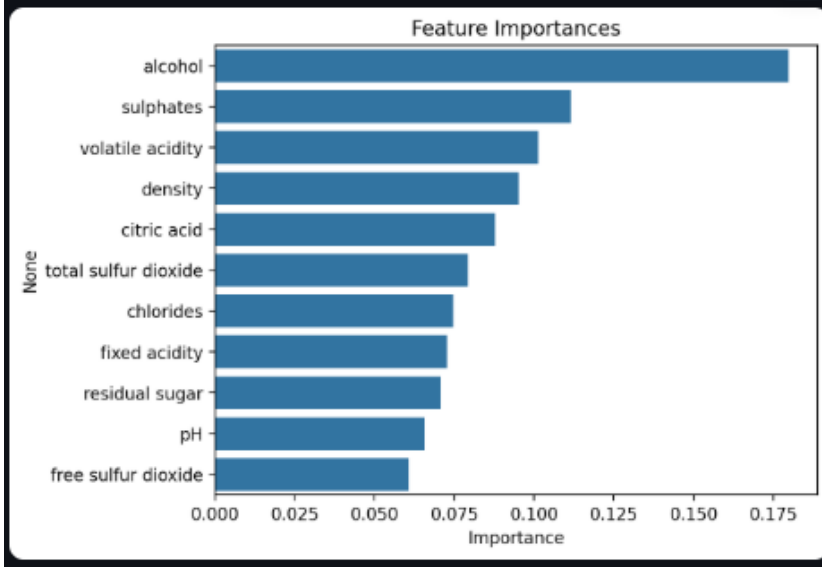
- Random Forest provides a measure of **feature importance**, which helps me explain which chemical properties (e.g., alcohol, sulphates, acidity) contribute most to wine quality. This interpretability is useful in the F&B industry.

6. Scalability and Ease of Use

- It is efficient to train on medium-sized datasets and easy to implement with libraries like **scikit-learn**.

Model Insights

This chart shows which properties are most important for predicting quality.



Dashboard to show results:

